



NeoFS: Practical introduction

NEO Saint Petersburg Competence Center
Alexey Vanin, Stanislav Bogatyrev

Neo Frontier Launchpad
June 2021

nspcc.ru



Neo St. Petersburg Competence Center

R&D company based in St. Petersburg, Russia

- Started July 27, 2018
- Deep expertise in data storage and distributed systems
- Close collaboration with universities
- Team of experienced researchers and engineers
- Maintenance of core Neo Blockchain infrastructure
- Neo Core development
- NeoFS, NeoGo and smaller ecosystem projects





NeoFS

Overall view


Decentralized object storage

InnerRing and Storage node types

Actions synced via Blockchain

gRPC API, support for S3 и HTTP

Storage Policy and ACL support





How does NeoFS store data?

Containers, Objects and their attributes

Address: ContainerID/ObjectID

ContainerID: Hash(Container)

Container: Storage Policy + Attributes

ObjectID: Hash(Header+Payload)

Object: Payload + Attributes

Container – like directory or bucket

Stores onchain (SideChain)

Object – like file

Stored offchain on NeoFS Storage Nodes

Address: 2UVmq...3PcdbSv / 8pw9Ma...HhEag5ds7AbDTF5BT



2Uwm...PcdbSv



8pw9...ds7AbDTF5BT

Container
OwnerID
BasicACL
Attributes
...
Placement Policy

Object
OwnerID
ContainerID
Attributes
...
Payload

Handling large objects

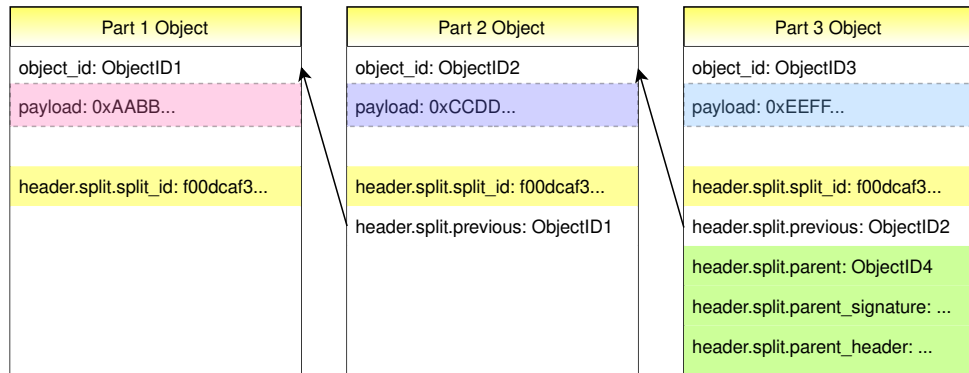
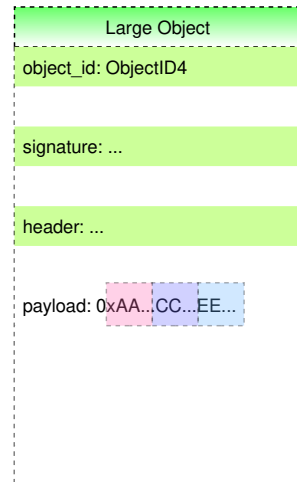
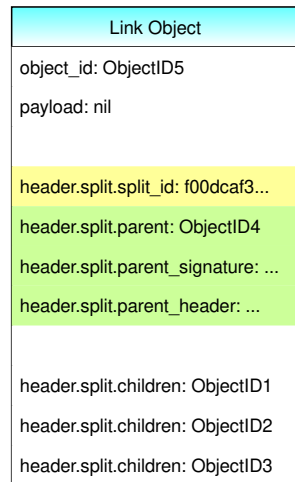
Transparent stream upload

Large objects are split automatically

The whole large object and each of it's parts are available

Stream upload/download with transparent object reconstruction

Compound objects are just objects =)



Network Map

Nodes and where to find them

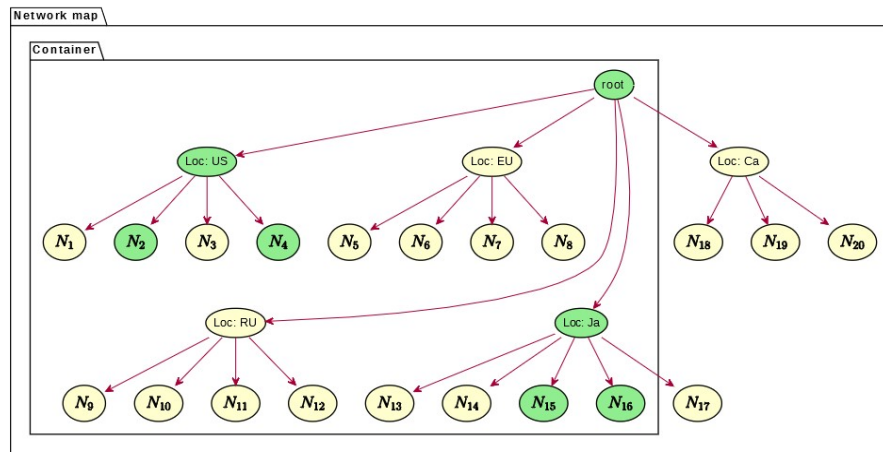
Network map contains active storage nodes

Storage node has a set of Attributes

Storage Policy works with attributes grouped into graph subtrees

HRW selection down to the depth set

New version of netmap is issued by Inner Ring each Epoch and stored in SideChain



REP 2 IN OURNODES

CBF 1

SELECT 4 FROM NSPCC AS OURNODES

FILTER "Deployed" EQ "NSPCC" AS NSPCC



Storage Policy

Simple rules

Storage Policy works with node attributes from Network Map

Flexible policy language with translation to SQL-like, JSON, Blockly, and more

Replicas distributed between node groups

Selectors define a node group from the network map

REP 1 in CORP

REP 1 in EXT

REP 1 in EMP

CBF 4

SELECT 4 Node IN CorpSRV as CORP

SELECT 2 Node IN GoodAliens as EXT

SELECT 6 Node IN EmpWKS as EMP

FILTER CorpProperty EQ ACMECorp AS IsACME

FILTER @IsACME AND NodeType EQ "SRV" AS CorpSRV

FILTER @IsACME AND NodeType EQ "WKS" AS EmpWKS

FILTER Rating GE 5 AS GoodAliens

Data Audit

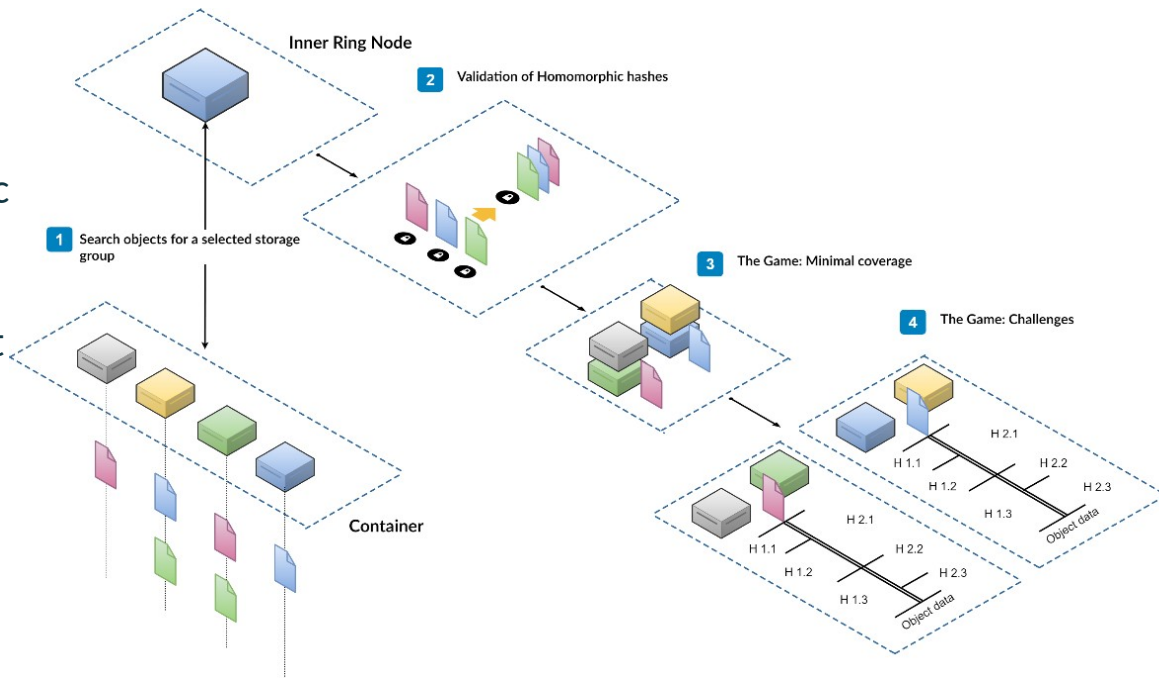
Multi stage game

Audit without data disclosure

ZK Proof based on homomorphic hashing

Works in untrusted environment

Traceable audit results onchain





Access Control List

BasicACL

Basic ACL for each container

Immutable

May be extended

Looks like POSIX file
Permissions =)

0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	2	2	2	2	2	2	2	2	2	3	3	3	
1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9	0	1	2
1				2				3				4				5				6				7				8			
0	0	0	1	1	1	0	0	1	0	0	0	1	1	0	0	1	0	0	0	1	1	0	0	1	1	0	0	1	1	0	0
				U	S	O	B	U	S	O	B	U	S	O	B	U	S	O	B	U	S	O	B	U	S	O	B	U	S	O	B
RSRV		X	F	GetRangeHas h				GetRange				Search				Delete				Put				Head				Get			



Access Control List

ExtendedACL

May be changed, but can't conflict with Basic ACL

Stored onchain in SideChain

Filters applied to both object and request fields and attributes

```
{
  "records": [
    {
      "operation": "GET",
      "action": "DENY",
      "filters": [
        {
          "headerType": "OBJECT",
          "matchType": "STRING_NOT_EQUAL",
          "key": "Classification",
          "value": "Public"
        }
      ],
      "targets": [
        {
          "role": "OTHERS"
        }
      ]
    }
  ]
}
```

Access Control List

BearerToken

May replace eACL for particular request

Just like JWT in modern Web

May limit the requesters

May be used via HTTP and other protocol gateways

Useful for complex authz cases

```
{
  "body": {
    "eacTable": {
      "version": {...},
      "containerID": {
        "value": "Ab1X...UeZTz8ZAVtzH5NWfKbC"
      },
      "records": [...],
      "ownerID": null,
      "lifetime": {
        "exp": "100500",
        "nbf": "2",
        "iat": "1"
      }
    },
    "signature": {
      "key": "A/ZjE3Ys1tvs6vVP7wPTxwqHYFlgoDhRnoeNBUwTq9Tq",
      "signature": "BCDU1mdffKmo9xVv/8H8XOmSQYN...uvtGNXTuDCWgHuSHZUcwqHtvXjsltBbl4gA0="
    }
  }
}
```



Practical demonstration


Get in touch with NeoFS

DEMO

NeoFS.Send

Web Application example

- Fully functional modern web app
- Serverless mode for real DApps
- Simple CDN example




Multi-level caching

GeoDNS

SmartContracts in control

Optimal data placement with NeoFS
Storage Policy

Local regulations support





Thank You!

Q&A

E-mail:

- *stanislav@nspcc.ru*
- *alexey@nspcc.ru*

Neo SPCC: <https://nspcc.ru>

NeoFS: <https://fs.neo.org>
<https://github.com/nspcc-dev/neofs-node>



NeoGo: <https://github.com/nspcc-dev/neo-go>

GitHub: <https://github.com/nspcc-dev/>

Medium: <https://medium.com/@neospcc/>


DHT and HashRing

Balance and data migration



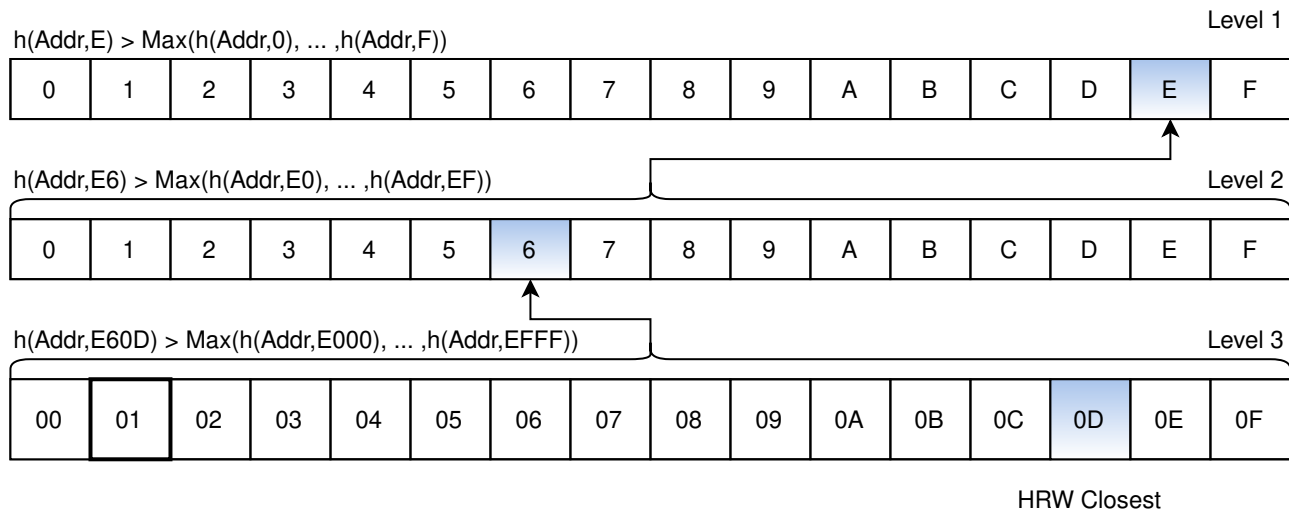
CRUSH in Ceph

Policy for the whole cluster



Rendezvous hashing

Calculating hash distance



Storage Engine

Rendezvous with Blobovnicza

1. Shards with HRW data placement
2. Local placement cache
3. Shallow dir structure placement
4. Hash-based placement in Blobovnicza
5. HRW-based placement in Blobovnicza

